

Time series forecasting

Initial setup

Import required packages

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

Import time series data: Airline passenger traffic

```
In [2]: data = pd.read_csv('airline-passenger-traffic.csv', header = None)
data.columns = ['Month', 'Passengers']
data['Month'] = pd.to_datetime(data['Month'], format='%Y-%m')
data = data.set_index('Month')
data.head(15)
```


```
Out [2]:
```

Passengers	
Month	
1949-01-01	112.0
1949-02-01	116.0
1949-03-01	132.0
1949-04-01	129.0
1949-05-01	121.0
1949-06-01	135.0
1949-07-01	148.0
1949-08-01	148.0
1949-09-01	136.0
1949-10-01	119.0
1949-11-01	104.0
1949-12-01	116.0

Time series analysis

Plot time series data

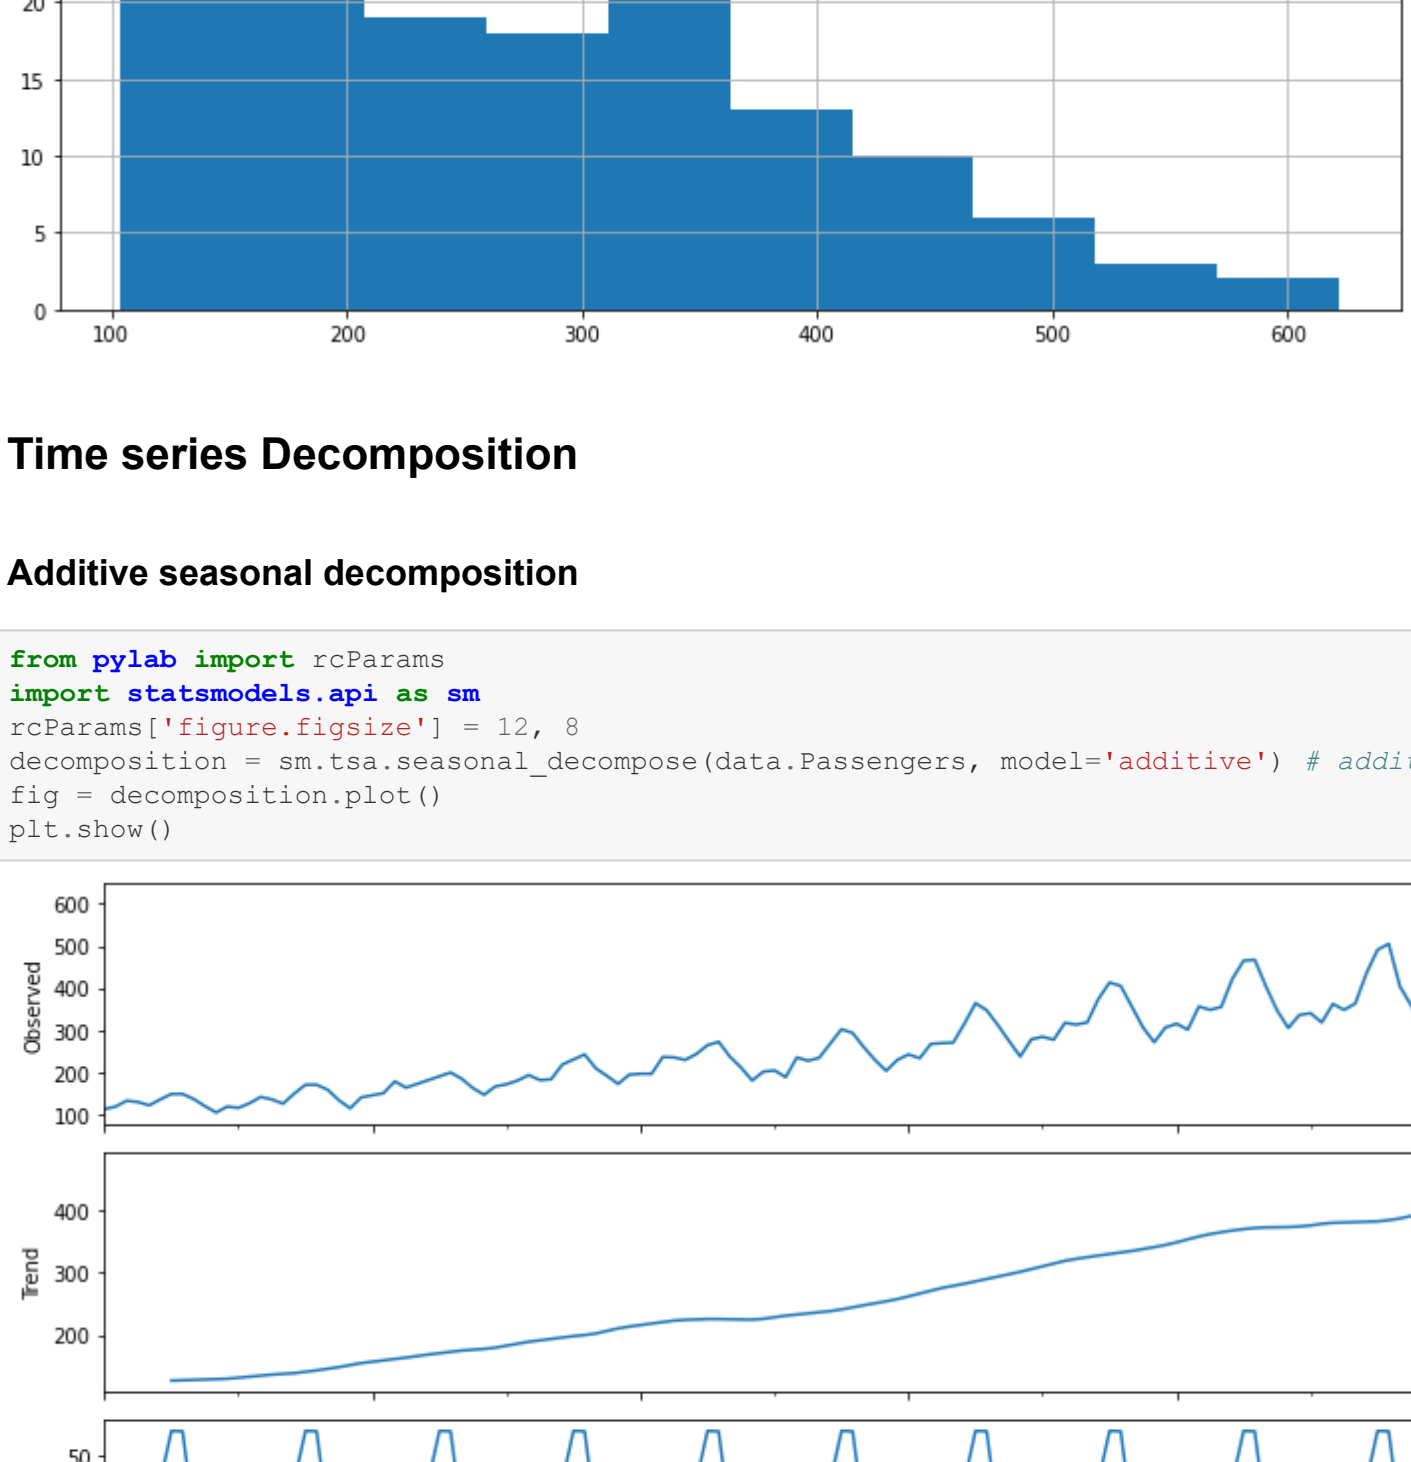
```
In [3]: data.plot(figsize=(12, 4))
plt.legend(loc='best')
plt.title('Airline passenger traffic')
plt.show(block=False)
```



Missing value treatment

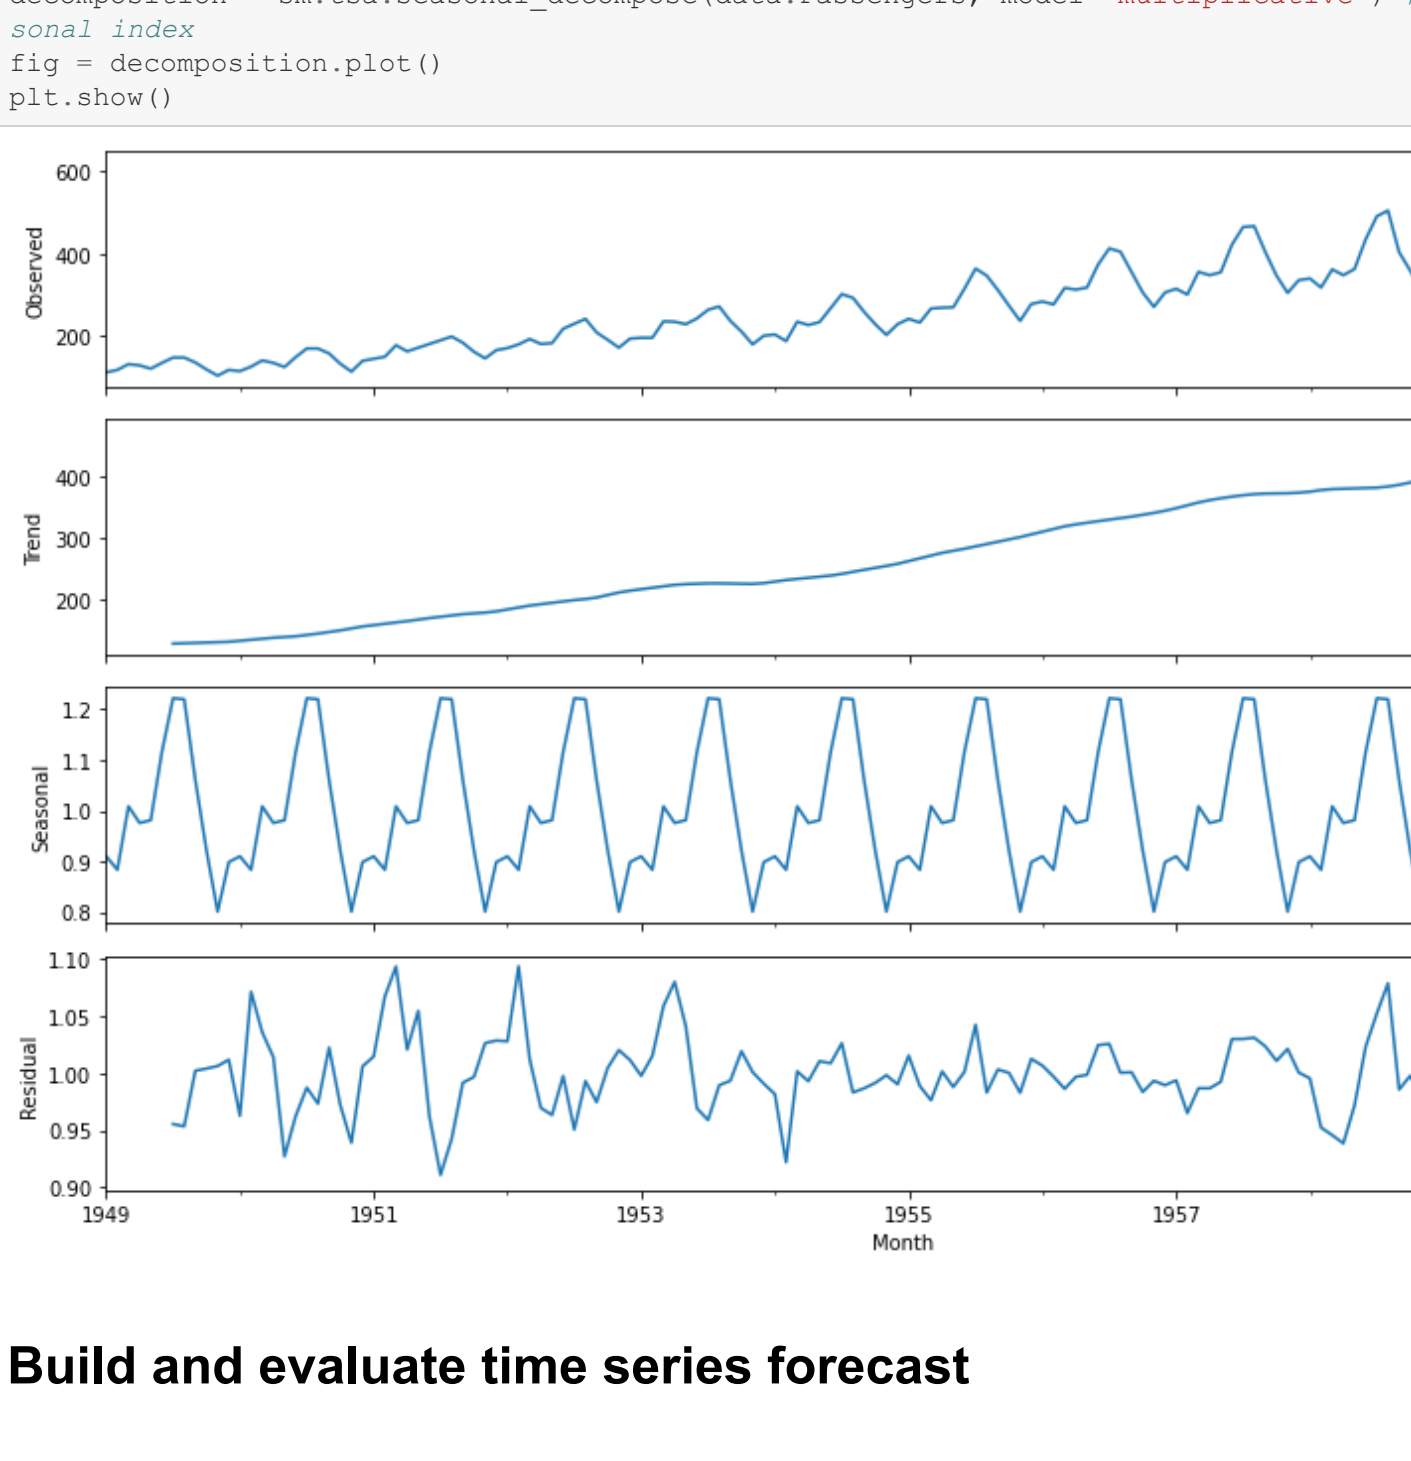
Mean imputation

```
In [4]: data = data.assign(Passengers_Mean_Imputation=data.Passengers.fillna(data.Passengers.mean()))
data[['Passengers_Mean_Imputation']].plot(figsize=(12, 4))
plt.legend(loc='best')
plt.title('Airline passenger traffic: Mean imputation')
plt.show(block=False)
```



Linear interpolation

```
In [5]: data = data.assign(Passengers_Linear_Interpolation=data.Passengers.interpolate(method='linear'))
data[['Passengers_Linear_Interpolation']].plot(figsize=(12, 4))
plt.legend(loc='best')
plt.title('Airline passenger traffic: Linear interpolation')
plt.show(block=False)
```




Use linear interpolation to impute missing values

```
In [6]: data['Passengers'] = data['Passengers_Linear_Interpolation']
data.drop(columns=['Passengers_Mean_Imputation', 'Passengers_Linear_Interpolation'], inplace=True)
```

Outlier detection

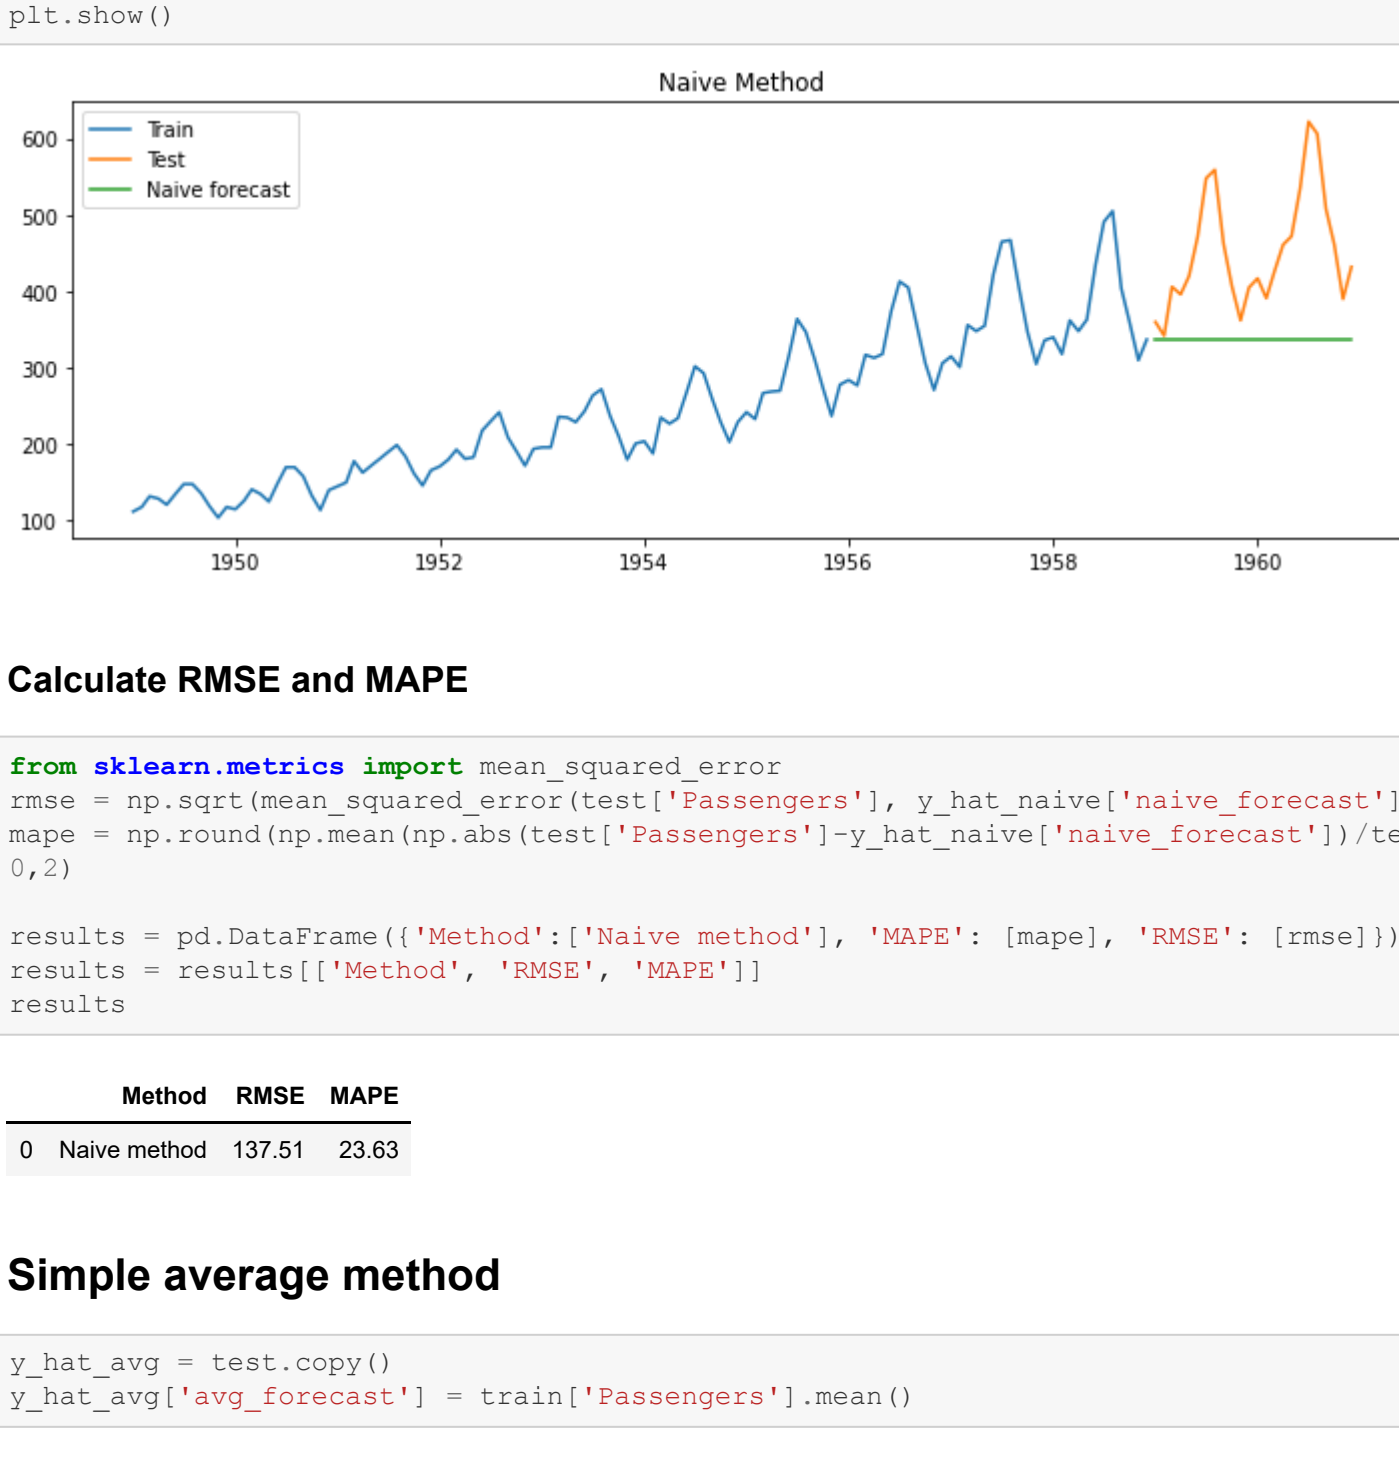
Box plot and interquartile range

```
In [7]: import seaborn as sns
fig = plt.subplots(figsize=(12, 2))
ax = sns.boxplot(x=data['Passengers'], whis=1.5)
```



Histogram plot

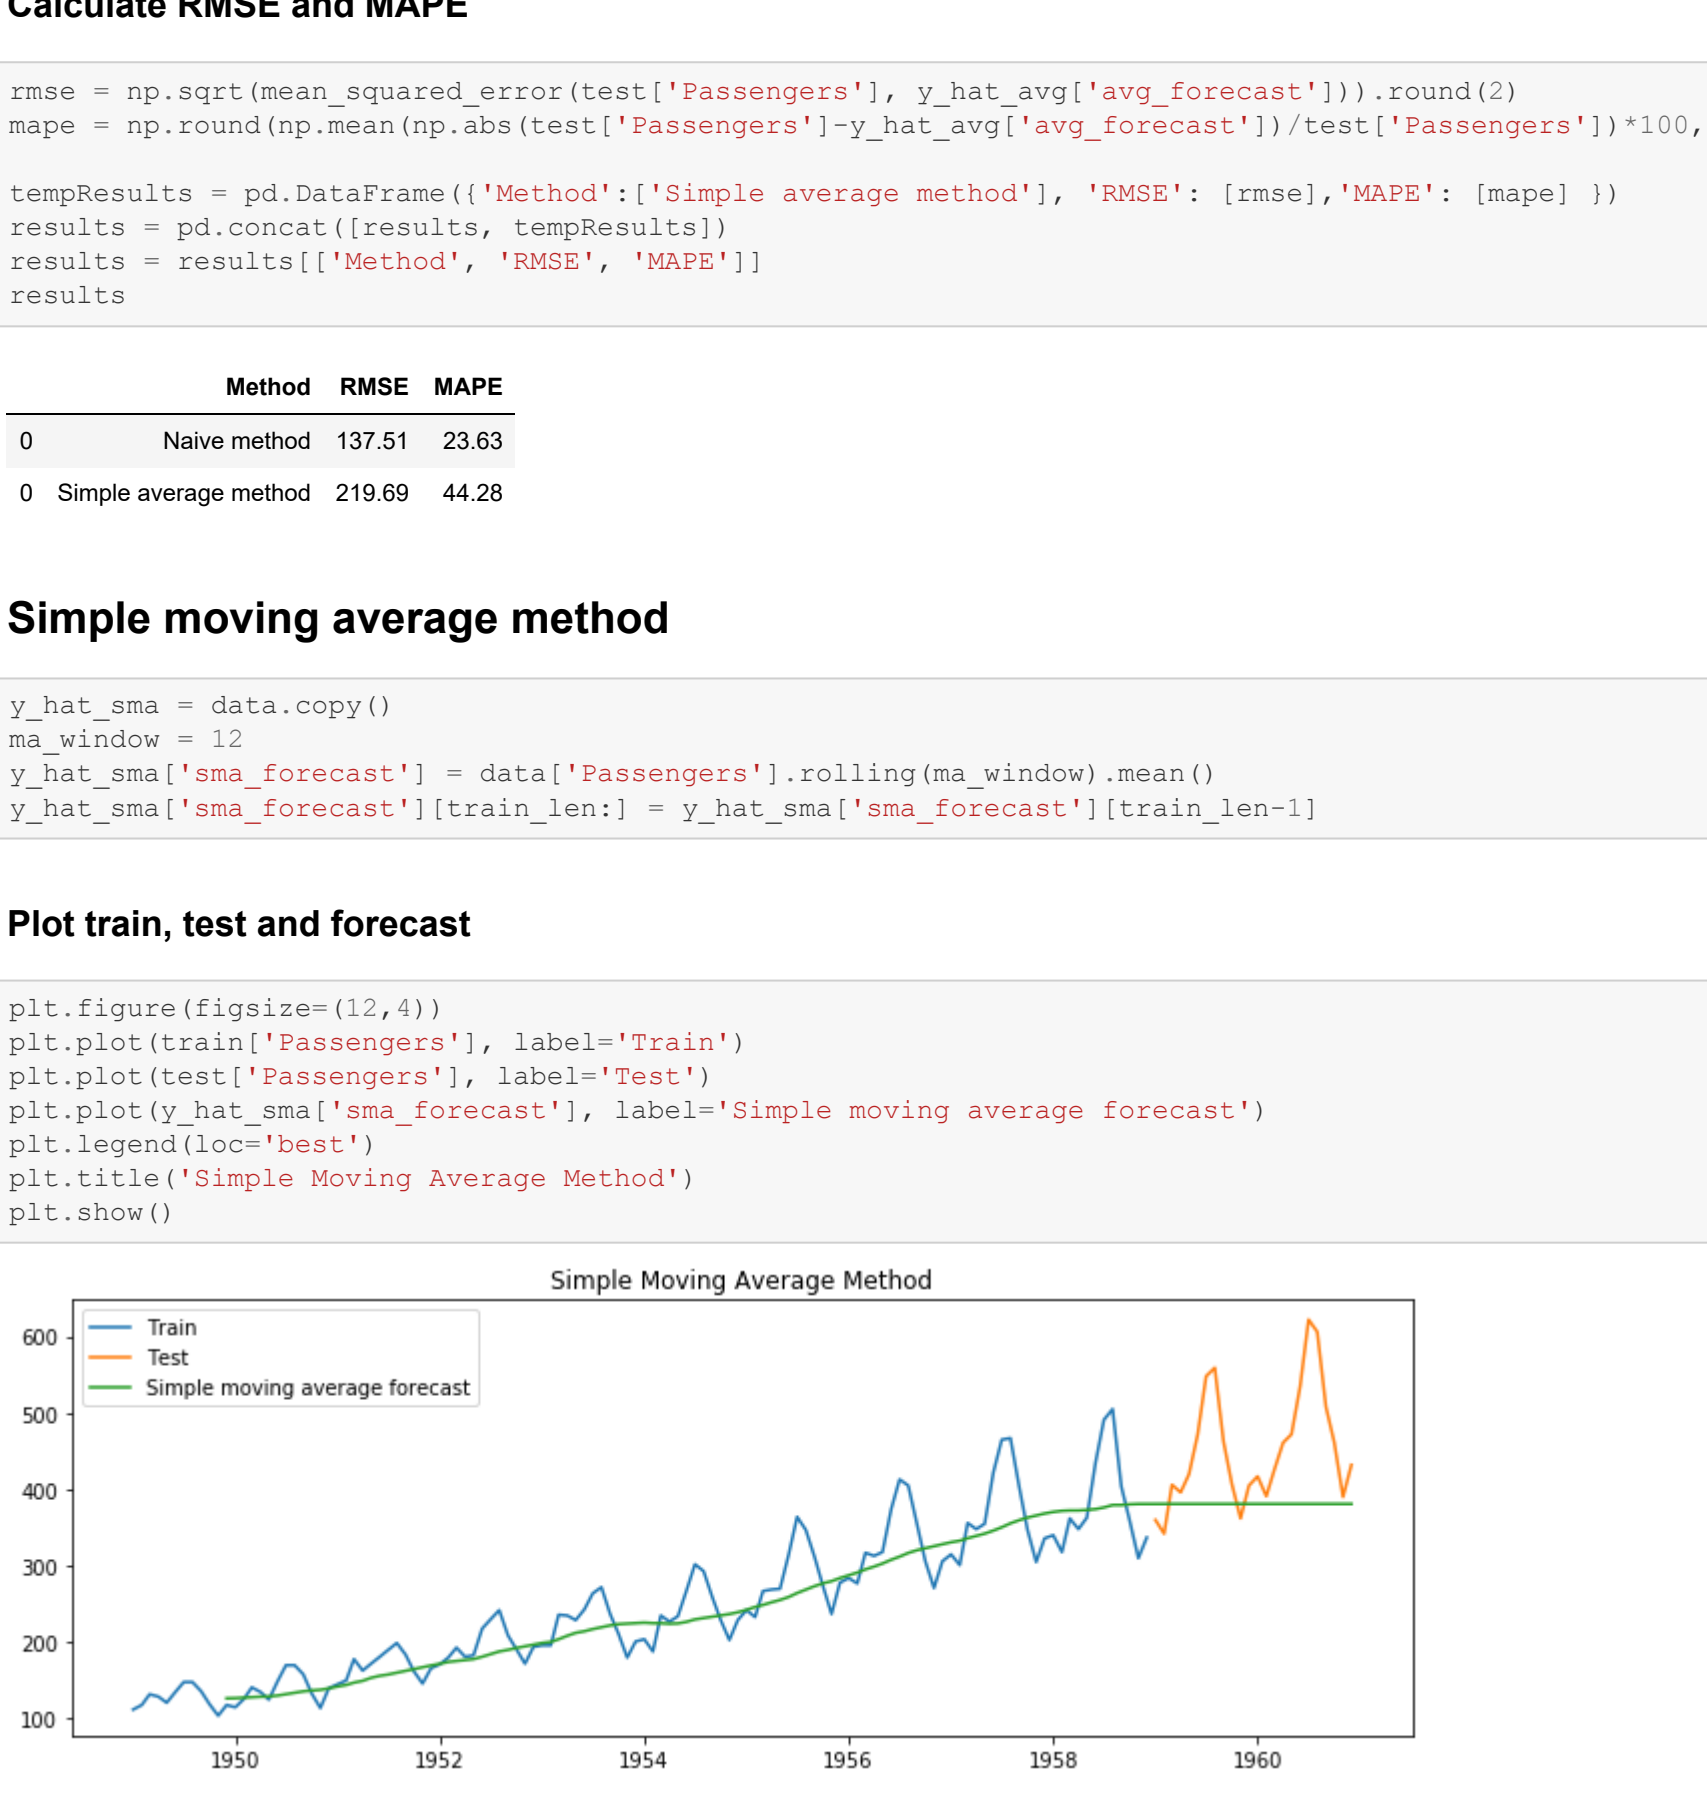
```
In [8]: fig = data.Passengers.hist(figsize = (12,4))
```



Time series Decomposition

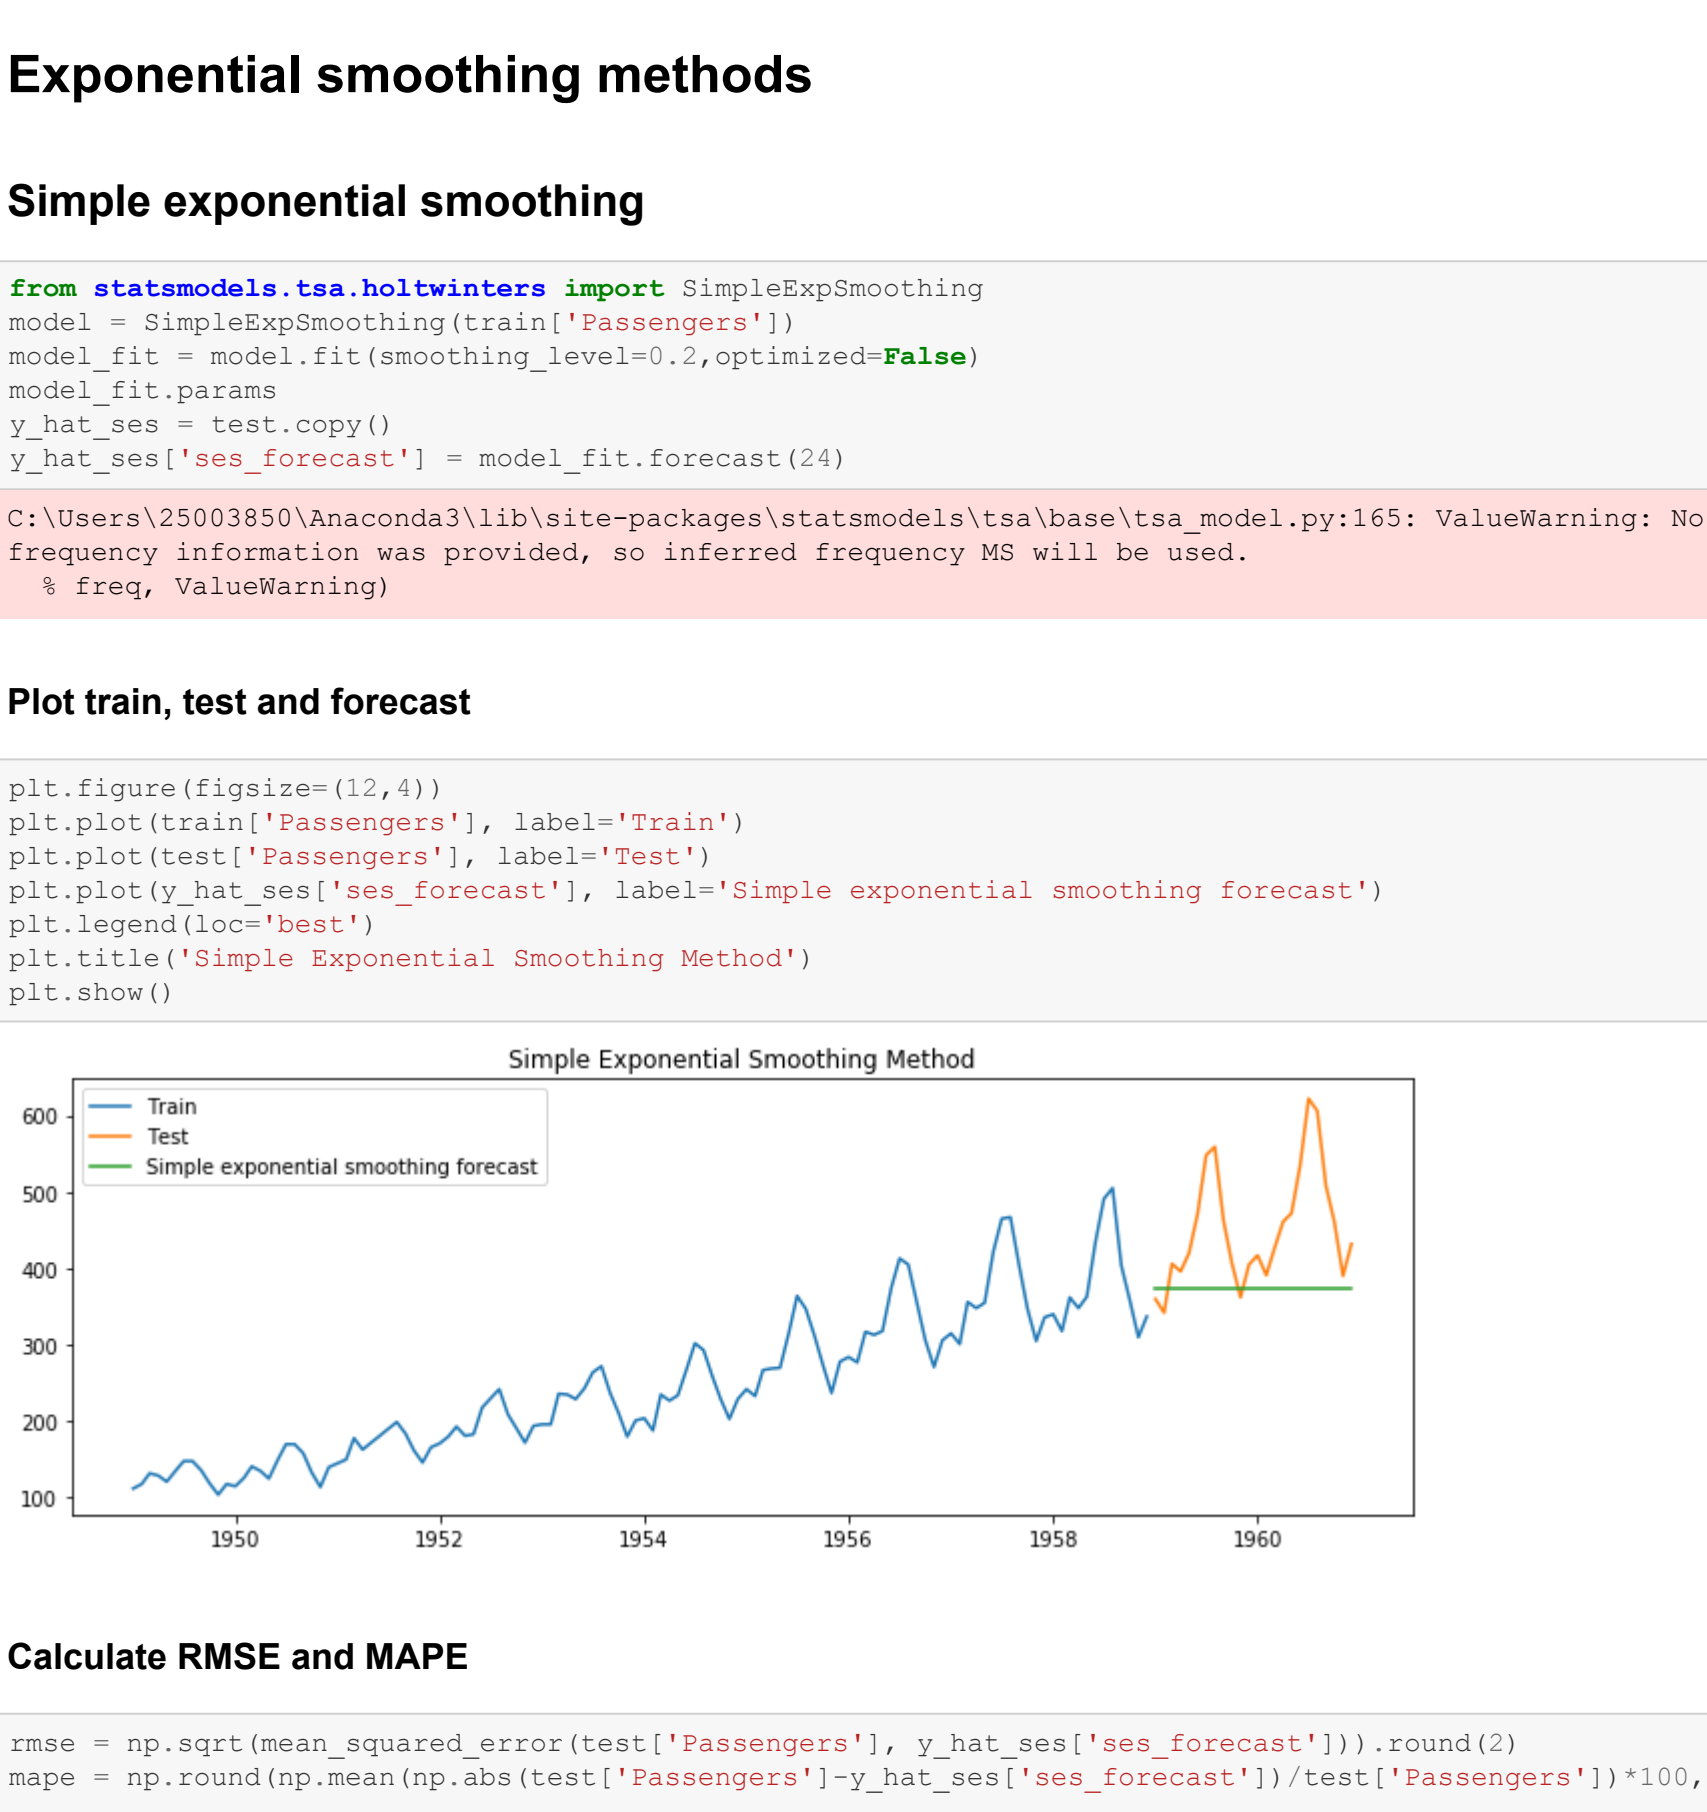
Additive seasonal decomposition

```
In [9]: from pylab import rcParams
import statsmodels.api as sm
rcParams['figure.figsize'] = 12, 8
decomposition = sm.tsa.seasonal_decompose(data.Passengers, model='additive') # additive seasonal index
fig = decomposition.plot()
plt.show()
```



Multiplicative seasonal decomposition

```
In [10]: decomposition = sm.tsa.seasonal_decompose(data.Passengers, model='multiplicative') # multiplicative sea
fig = decomposition.plot()
plt.show()
```



Build and evaluate time series forecast

Split time series data into training and test set

```
In [11]: train_len = 120
test = data[:train_len] # first 120 months as training set
test = data[train_len:] # last 24 months as out-of-time test set
```

Simple time series methods

Naive method

```
In [12]: y_hat_naive = test.copy()
y_hat_naive['naive_forecast'] = train['Passengers'][train_len-1]
```

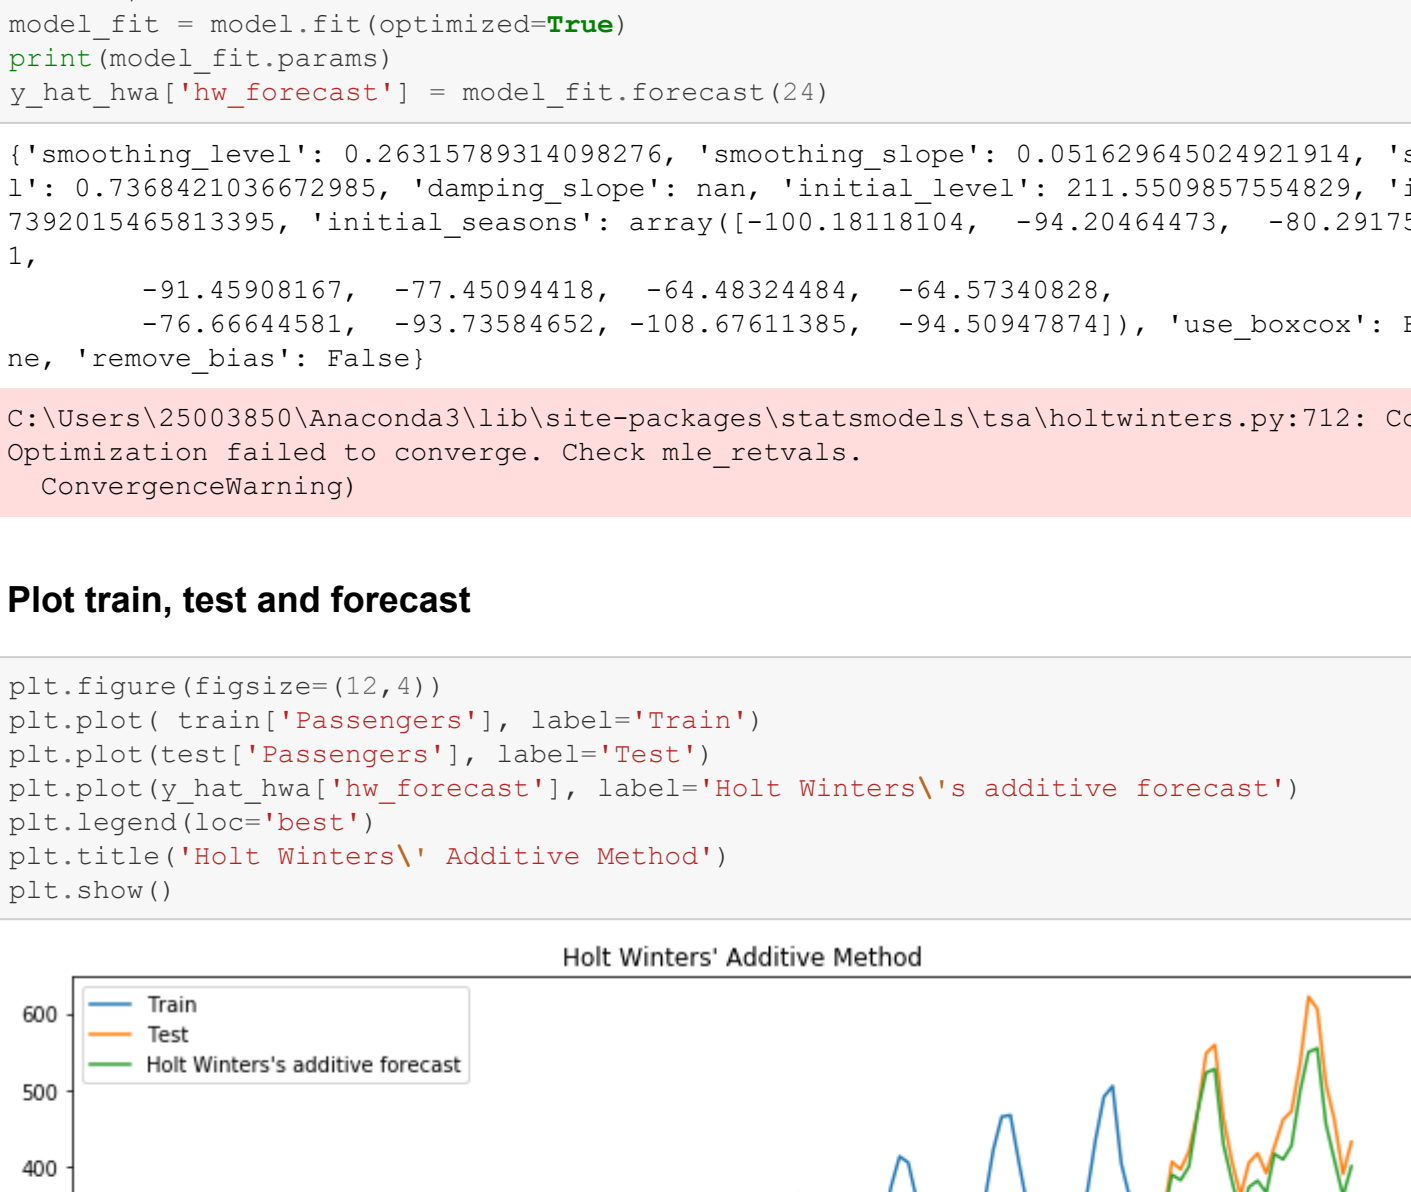
```
In [14]: y_hat_naive
```

```
Out [14]:
```

Month	Passengers	naive_forecast
1959-01-01	360.0	337.0
1959-02-01	342.0	337.0
1959-03-01	406.0	337.0
1959-04-01	396.0	337.0
1959-05-01	420.0	337.0
1959-06-01	472.0	337.0
1959-07-01	548.0	337.0
1959-08-01	559.0	337.0
1959-09-01	483.0	337.0
1959-10-01	407.0	337.0
1959-11-01	362.0	337.0
1959-12-01	405.0	337.0
1960-01-01	417.0	337.0
1960-02-01	391.0	337.0
1960-03-01	426.0	337.0
1960-04-01	461.0	337.0
1960-05-01	473.0	337.0
1960-06-01	535.0	337.0
1960-07-01	622.0	337.0
1960-08-01	606.0	337.0
1960-09-01	508.0	337.0
1960-10-01	461.0	337.0
1960-11-01	390.0	337.0
1960-12-01	432.0	337.0

Plot train, test and forecast

```
In [13]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_naive['naive_forecast'], label='Naive forecast')
plt.legend(loc='best')
plt.title('Naive Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [14]: from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_naive['naive_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_naive['naive_forecast']))/test['Passengers'])*100,2)

results = pd.DataFrame({'Method': ['Naive method'], 'RMSE': [rmse], 'MAPE': [mape]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [14]:
```

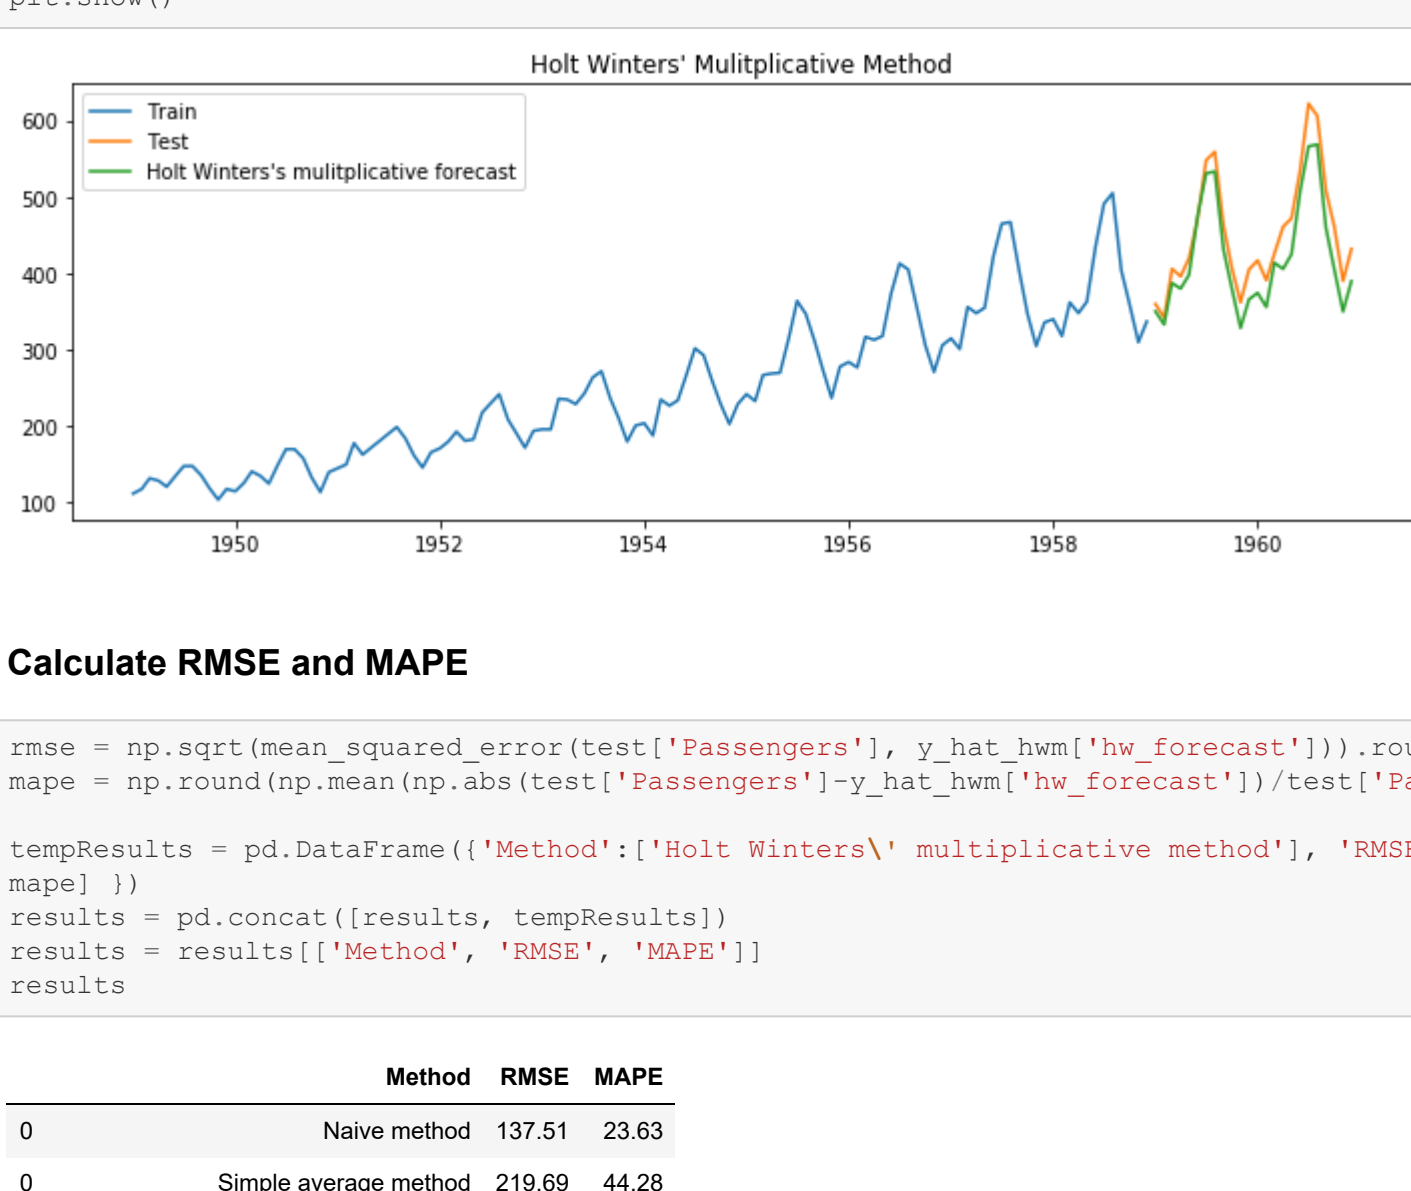
	Method	RMSE	MAPE
0	Naive method	137.51	23.63

Simple average method

```
In [15]: y_hat_avg = test.copy()
y_hat_avg['avg_forecast'] = train['Passengers'].mean()
```

Plot train, test and forecast

```
In [16]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_avg['avg_forecast'], label='Simple average forecast')
plt.legend(loc='best')
plt.title('Simple Average Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [17]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_avg['avg_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_avg['avg_forecast']))/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Simple average method'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [17]:
```

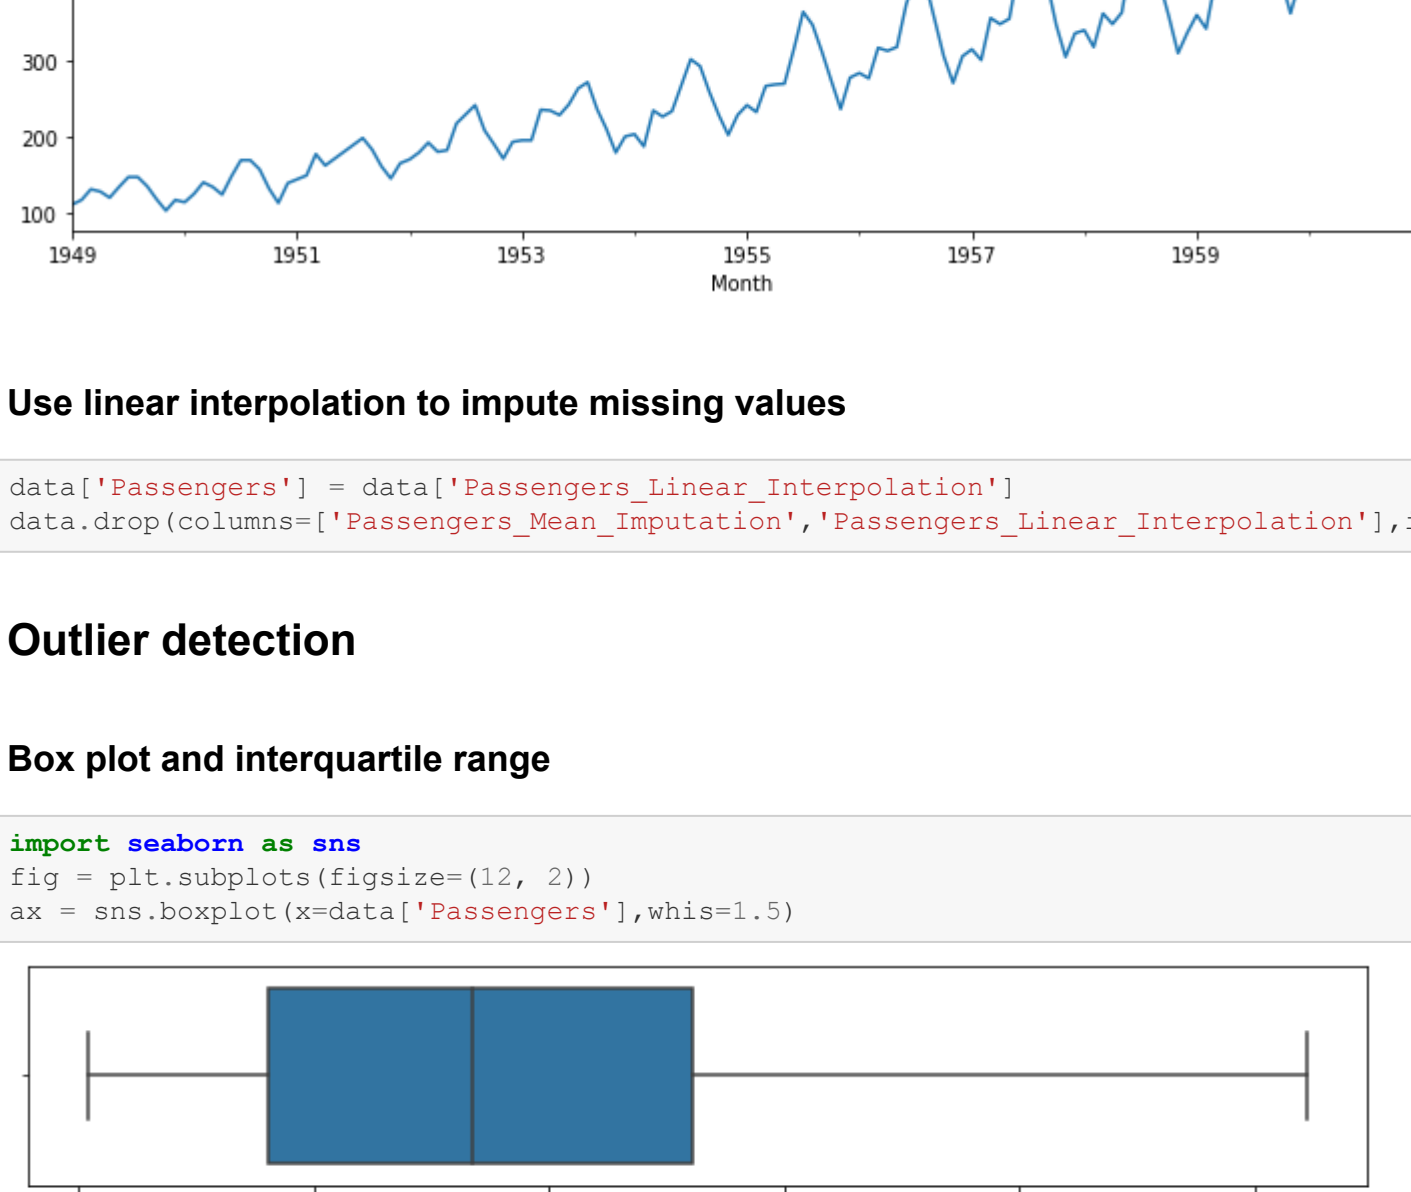
	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28

Simple moving average method

```
In [18]: y_hat_sma = test.copy()
ma_window = 12
y_hat_sma['sma_forecast'] = data['Passengers'].rolling(ma_window).mean()
y_hat_sma['sma_forecast'][train_len:] = y_hat_sma['sma_forecast'][train_len-1]
```

Plot train, test and forecast

```
In [19]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_sma['sma_forecast'], label='Simple moving average forecast')
plt.legend(loc='best')
plt.title('Simple Moving Average Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [20]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_sma['sma_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_sma['sma_forecast'])/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Simple moving average forecast'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [20]:
```

	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28
0	Simple moving average forecast	103.33	15.54

Exponential smoothing methods


Simple exponential smoothing

```
In [21]: from statsmodels.tsa.holtwinters import SimpleExpSmoothing
model = SimpleExpSmoothing(train['Passengers'])
model_fit = model.fit(smoothing_level=0.2, optimized=False)
model_fit.params
y_hat_hwm['hwm_forecast'] = model_fit.forecast(24)
```

```
C:\Users\25003850\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:165: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
%Freq, ValueWarning)
```

Plot train, test and forecast

```
In [22]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_hwm['hwm_forecast'], label='Simple exponential smoothing forecast')
plt.legend(loc='best')
plt.title('Simple Exponential Smoothing Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [23]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_hwm['hwm_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_hwm['hwm_forecast'])/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Simple exponential smoothing forecast'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [23]:
```

	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28
0	Simple moving average forecast	103.33	15.54
0	Simple exponential smoothing forecast	107.65	16.49
0	Simple exponential smoothing method	107.65	16.49

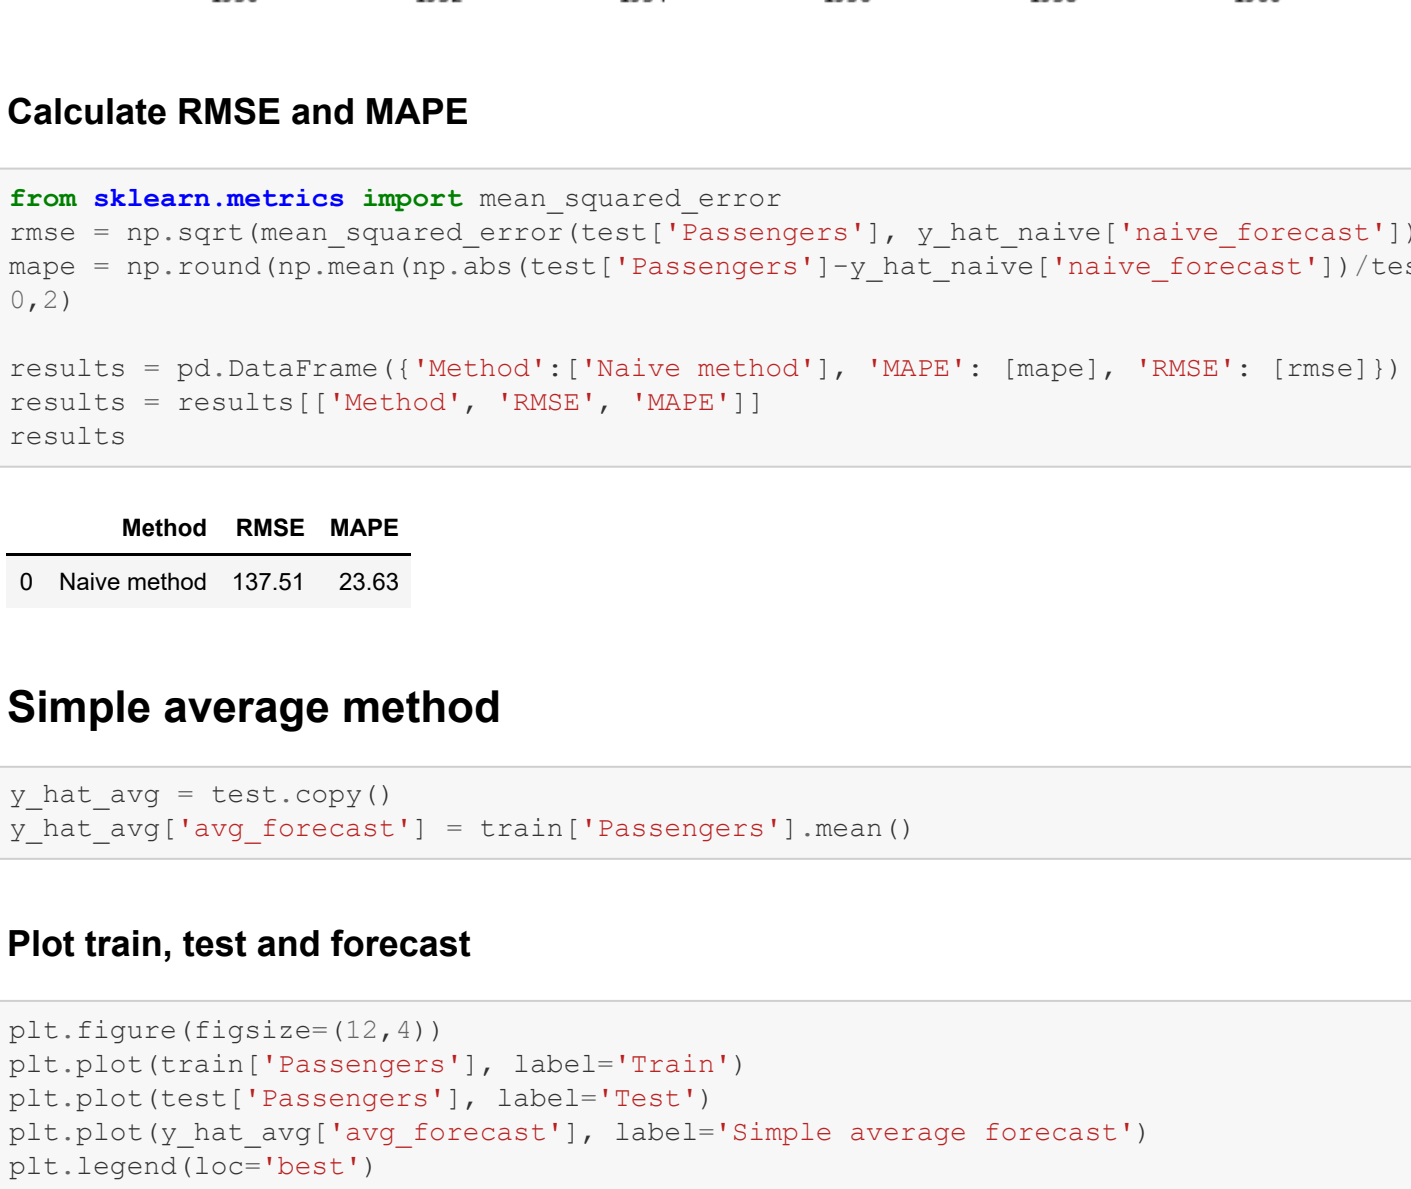
Holt's method with trend

```
In [24]: from statsmodels.tsa.holtwinters import ExponentialSmoothing
model = ExponentialSmoothing(np.asarray(train['Passengers']), seasonal_periods=12, trend='additive', se
seasonal=None)
model_fit = model.fit(smoothing_level=0.2, smoothing_slope=0.01, optimized=False)
print(model_fit.params)
y_hat_hwm['hwm_forecast'] = model_fit.forecast(len(test))
```

```
{'smoothing_level': 0.2, 'smoothing_slope': 0.01, 'smoothing_seasonal': None, 'damping_slope': nan, 'initial_level': 112.0, 'initial_slope': 6.0, 'initial_seasons': array([], dtype=float64), 'use_boxco
x': False, 'lambda': None, 'remove_bias': False}
```

Plot train, test and forecast

```
In [25]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_holt['holt_forecast'], label='Holt's exponential smoothing forecast')
plt.legend(loc='best')
plt.title('Holt's Exponential Smoothing Method')
plt.show()
```



Calculate RSME and MAPE

```
In [26]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_holt['holt_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_holt['holt_forecast'])/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Holt's exponential smoothing method'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [26]:
```

	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28
0	Simple moving average forecast	103.33	15.54
0	Simple exponential smoothing forecast	107.65	16.49
0	Holt's exponential smoothing method	107.65	16.49
0	Holt's exponential smoothing method	71.94	11.11

Holt Winter's additive method with trend and seasonality

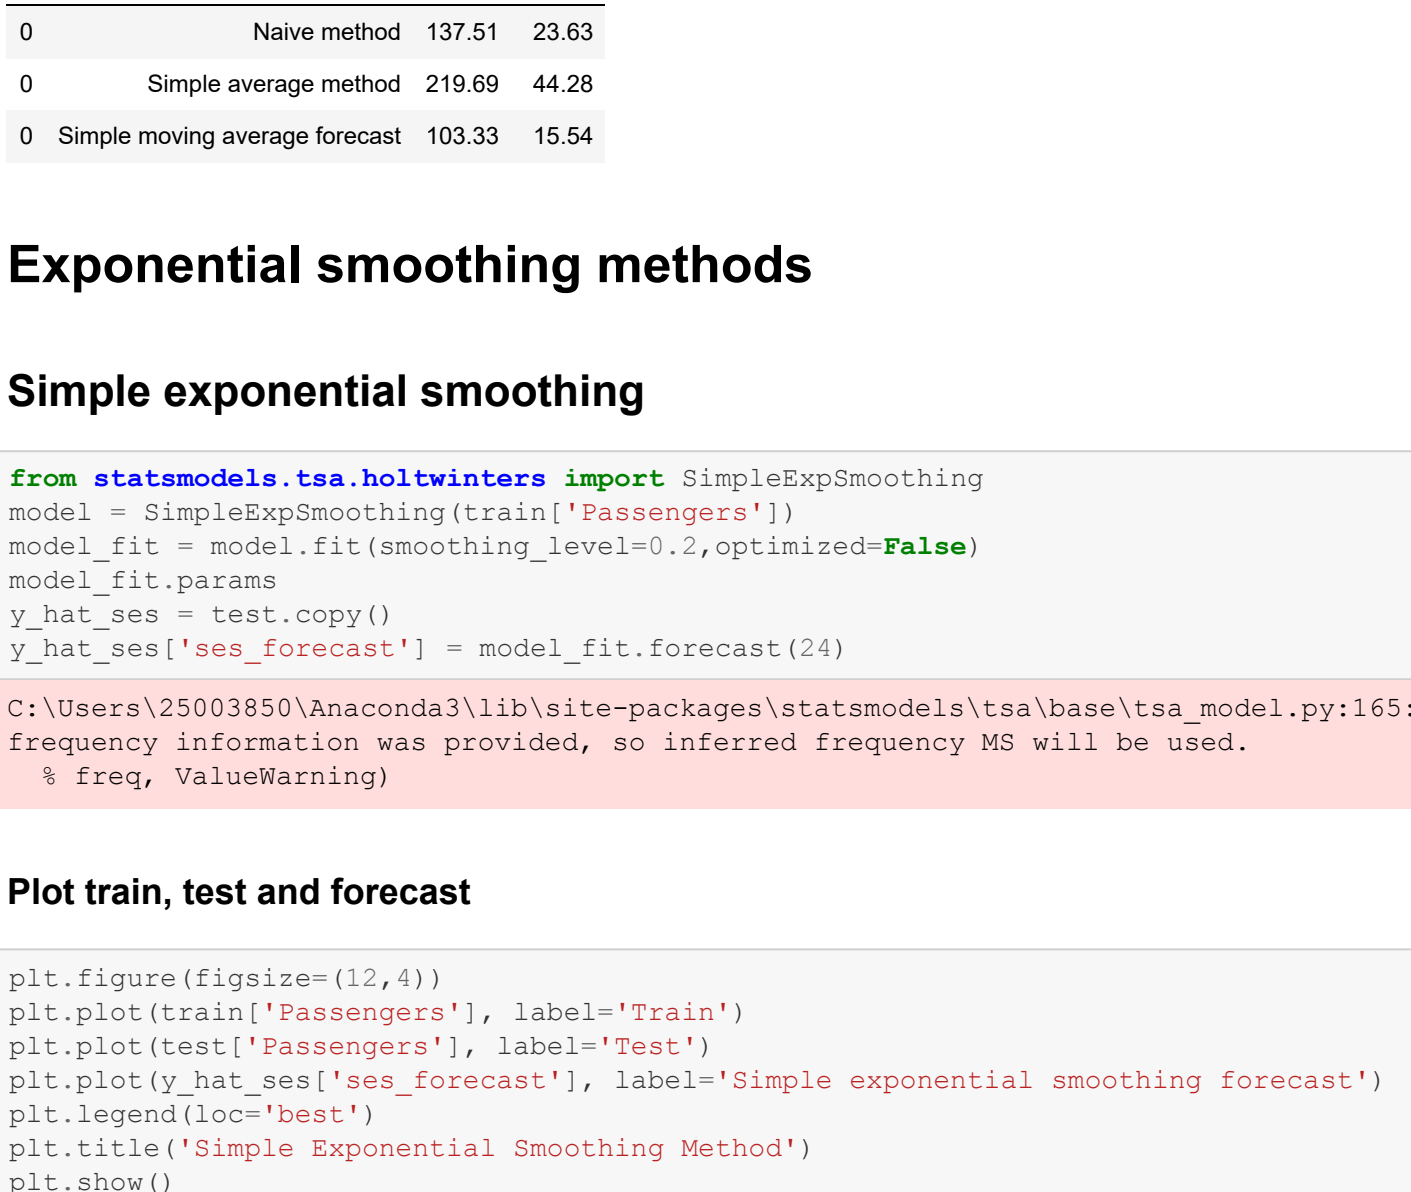
```
In [27]: y_hat_hwa = test.copy()
model = ExponentialSmoothing(np.asarray(train['Passengers']), seasonal_periods=12, trend='add', seasona
l='mul')
model_fit = model.fit(optimized=True)
print(model_fit.params)
y_hat_hwa['hwa_forecast'] = model_fit.forecast(24)
```

```
{'smoothing_level': 0.26315789314098276, 'smoothing_slope': 0.051629645024921914, 'smoothing_seasonal': 0.61599393691, 'damping_slope': nan, 'initial_level': 212.39830870038878, 'initial_slope': 1.105
1.7392015456813395, 'initial_seasons': array([-100.1818104, -94.20464473, -80.29175401, -63.4175839
1, -49.45908167, -37.45094418, -24.48324484, -11.50947874]), 'use_boxcox': False, 'lambda': No
ne, 'remove_bias': False}
```

```
C:\Users\25003850\Anaconda3\lib\site-packages\statsmodels\tsa\hwtwinters.py:712: ConvergenceWarning: Optimization failed to converge. Check mle_retvals
ConvergenceWarning)
```

Plot train, test and forecast

```
In [28]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_hwa['hwa_forecast'], label='Holt Winters's additive forecast')
plt.legend(loc='best')
plt.title('Holt Winters' Additive Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [29]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_hwm['hwm_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_hwm['hwm_forecast'])/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Holt Winters\' multiplicative method'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [29]:
```

	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28
0	Simple moving average forecast	103.33	15.54
0	Simple exponential smoothing forecast	107.65	16.49
0	Holt's exponential smoothing method	71.94	11.11
0	Holt Winters' additive method	35.10	6.53

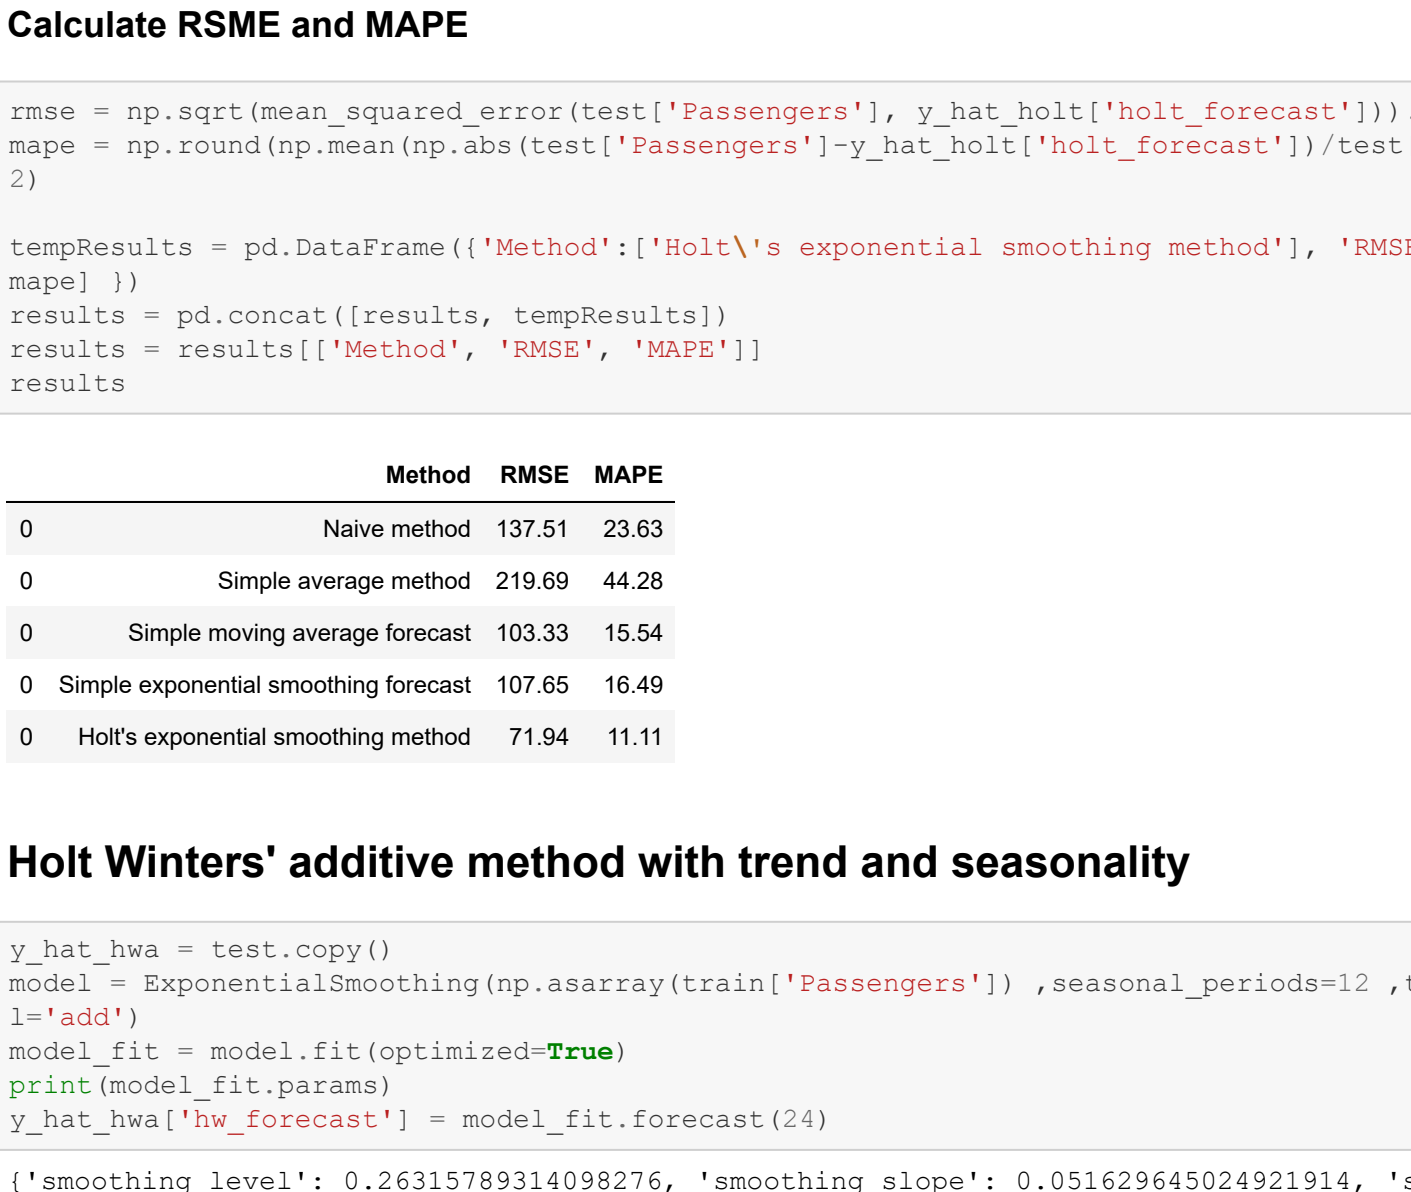
Holt Winter's multiplicative method with trend and seasonality

```
In [30]: y_hat_hwm = test.copy()
model = ExponentialSmoothing(np.asarray(train['Passengers']), seasonal_periods=12, trend='add', seasona
l='mul')
model_fit = model.fit(optimized=True)
print(model_fit.params)
y_hat_hwm['hwm_forecast'] = model_fit.forecast(24)
```

```
{'smoothing_level': 0.38484064840698556, 'smoothing_slope': 0.035386453247783, 'smoothing_seasonal': 0.61599393691, 'damping_slope': nan, 'initial_level': 212.39830870038878, 'initial_slope': 1.105
0.53222879, 'initial_seasons': array([0.51755725, 0.54595018, 0.60261532, 0.56966257, 0.52017629,
0.43937027, 0.31362925]), 'use_boxcox': False, 'lambda': None, 'remove_bias': False}
```

Plot train, test and forecast

```
In [31]: plt.figure(figsize=(12,4))
plt.plot(train['Passengers'], label='Train')
plt.plot(test['Passengers'], label='Test')
plt.plot(y_hat_hwm['hwm_forecast'], label='Holt Winters's multiplicative forecast')
plt.legend(loc='best')
plt.title('Holt Winters' Multiplicative Method')
plt.show()
```



Calculate RMSE and MAPE

```
In [32]: rmse = np.sqrt(mean_squared_error(test['Passengers'], y_hat_hwm['hwm_forecast'])).round(2)
mape = np.round(np.mean(np.abs(test['Passengers']-y_hat_hwm['hwm_forecast'])/test['Passengers'])*100,2)

tempResults = pd.DataFrame({'Method': ['Holt Winters\' multiplicative method'], 'RMSE': [rmse], 'MAPE': [mape]})
results = pd.concat([results, tempResults])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
Out [32]:
```

	Method	RMSE	MAPE
0	Naive method	137.51	23.63
0	Simple average method	219.69	44.28
0	Simple moving average forecast	103.33	15.54
0	Simple exponential smoothing forecast	107.65	16.49
0	Holt's exponential smoothing method	71.94	11.11
0	Holt Winters' additive method	34.83	6.91